

Not So Slowth: Invertible VDF for Ethereum and others

Dmitry Khovratovich
Ethereum Foundation

Mary Maller
Ethereum Foundation

Pratyush Ranjan Tiwari
Ethereum Foundation

June 2, 2021

Contents

1	Introduction	2
2	Requirements	2
2.1	Delay statements	3
2.2	Delay hardware	3
2.3	Delay proofs	3
2.4	Cryptographic primitives and assumptions	4
2.5	Prover hardware system	4
3	VDF definitions	4
3.1	Generic VDF	4
3.2	VDF proofs based on Verifiable Computation	5
3.3	Invertible VDF	5
3.4	Invertible VDF flow	6
3.5	Sequentiality	6
4	Attacks on invertible VDFs	7
4.1	Precomputation attacks	7
4.1.1	Attack 1	7
4.1.2	Attack 2	7
4.1.3	Attack 3	8
4.1.4	Mitigation and parameter selection	8
4.2	Grobner Basis and Algebraic	8
5	Quantum algorithms	9
5.1	Grover's algorithm	9
5.1.1	Single-target Grover's	9
5.1.2	Multi-target Grover's	10
5.1.3	Quantum Parallelism	10
5.2	Attacks based on Grover's	10
5.2.1	Grover Attacks 1 and 2	10
5.2.2	Attacks that do not work	11
5.3	Practicality of Quantum Search	11
5.4	Quantum Grobner	11

6	Other attacks	12
6.1	MIMC reduction	12
6.2	Attacks on Verifiable Computation	12
7	Concrete Proposals	12
7.1	Recommendations for ϕ	12
7.2	MinRoot (Khovratovich-Feist)	13
7.3	VeeDo (StarkWare)	13
7.4	Sloth and Sloth++ (Wesolowski-Boneh et al.)	14
7.5	MinRoot Lite	15
8	Concrete security of MinRoot(s)	15
8.1	Classical attacks	15
8.2	Quantum attacks	16
9	Concrete security of other proposals	16
10	Performance	16
11	Conclusion	17

1 Introduction

Many distributed protocols benefit not only from fast but also from (guaranteed) slow computations. The latter is often required when a party is supposed to produce an output that will eventually benefit some other participants, so that the slowness of computation would guarantee unpredictability and thus fairness.

An example of this requirement in Ethereum 2.0 is the following. A group of 32 validators progressively builds a chain of collective randomness, also called a beacon chain, with one value O generated per epoch \mathcal{E} . The value O must be computed by the leading validator of \mathcal{E} and must depend on certain inputs produced by other validators during \mathcal{E} . When generated, O serves as the source of entropy in other chains and thus affect a number of computations potentially beneficial to the validator. It is thus of utmost importance to ensure that O can be outputted only after a certain *delay*. In other words, O is derived from a *verifiable delay function* (VDF).

A number of VDF constructions has been proposed in the recent past [LW17, Wes19, BBBF18, Pie19] with their own advantages and limitations. The VDF Alliance seeks for a concretely-efficient post-quantum VDF which may be used by L1 projects such as Cosmos, Ethereum, Filecoin, Tezos as well as L2 applications. The post-quantumness requirement rules out the promising RSA VDF [Wes19, Pie19], but leaves open the possibility for VDFs based on verifiable computation, which are the subject of this report.

2 Requirements

On the high level, a protocol party, when outputting value O , also proves a statement of kind

I have spent at least t time to produce O .

It is not important how exactly O is computed but it *is* important that the parameter t be close, or at least lower bound the actual time spent. The applications of VDF also yield separate requirements for hardware (where it is implemented), security properties, and statement parameters, which are all listed below.

2.1 Delay statements

A delay statement is a statement of the form

`delay(input, difficulty) == output`

or equivalently

$$F(I, d) = O$$

where I and O are integer (tuples) and d is a non-negative integer. The difficulty parameter regulates the execution time. We need certain delay granularity, i.e. the difficulty can vary with precision in seconds. This naturally suggests that F is an iteration of some function f with sufficient granularity. A single call to f is then declared the unit of difficulty d , with $d \approx 2^{24}$ (for proposals listed in Section 7) corresponding to 1 second. We restrict ourselves to delay statements with:

- 15 sec to 3 days **delay range** ($2^{28} \leq d \leq 2^{42}$);
- 1 sec **delay granularity** – ($d = 0 \pmod{2^{24}}$).

2.2 Delay hardware

The delay hardware is supposed to be a USB stick fitted with an ASIC implementing the VDF with an additional requirement: every second an intermediate output is streamed to the host in real time. As the hardware is not yet manufactured, it can be simulated by precomputing every about 2^{24} -th round outputs and reading the appropriate output every $2^{24} \cdot 64$ ns ($\approx 1s$).

2.3 Delay proofs

Existing VDF constructions rely on a separate proof system that proves correctness of the VDF computation and thus ensures the delay statement. Clearly, the security of VDF relies on the one of the proof system. The following requirements are given:

1. First, a proof system should provide reasonable security level in a widely accepted security model. A random oracle model is chosen for its simplicity and versatility. There are two flavours of it: *programmable* and *non-programmable ROM*. The former assumes we can dynamically select RO outputs in a simulation for the security proof, which simplifies the proof but formally contrasts with real-world instantiations of RO with hash functions, whose functionality is pre-fixed. The latter model is somewhat closer to the reality but such security proofs are not always available. We thus *prefer* the non-programmable ROM but it is not a strong requirement. Regarding the security level, at least 120 bits of security is required (possibly under additional assumptions on some primitives used in the system).
2. Secondly, a proof of a delay statement, whatever it is, should be *post-quantum secure*. Concretely, no quantum attacks should exist on underlying primitives, and additionally the security in the random oracle model, if being employed, should translate to quantum oracles. More precisely, at least 80 bits of **provable soundness in a quantum random oracle model** is required, with slight preference for non-programmable one.
3. The proof should be **compact**. It is understood that the post-quantum security makes it difficult to be truly succinct (e.g. a few kilobytes), but it should be possible to contain the proof into 200 kBytes.
4. The proof should be verified within very short time (**succinct verification**). Concretely, we require at most 20ms verification time on a Raspberry Pi 4 Model B (single thread, no overclocking).

2.4 Cryptographic primitives and assumptions

It is natural to accept widely used cryptographic primitives as secure ones. For instance standard hash functions such as SHA-256 and SHA-3 (Keccak) may be assumed to be ideal 256-bit hash function and may be used to instantiate a 256-bit random oracle in the Fiat-Shamir transformation and similar operations.

It should be noted however that the regular hash functions are not well suited for most proof systems capable of proving delay statements. As a result, their usage in delay proofs is probably limited.

2.5 Prover hardware system

We have a separate requirement to a potential prover hardware system. We expect it to be a prover rig with prover software producing delay proofs for $1 \leq u \leq 8$ always-active delay hardware units. The requirement is to design and implement a prover system with

- **commodity hardware** – the prover rig is built from off-the-shelf commodity hardware and requires little assembly;
- **low latency** – delay proofs are produced with a latency (relative to the delay function output produced by the delay hardware) of:
 - at most 10% of the delay time if $d \geq 2^{36}$ (i.e. more than an hour);
 - at most 20% of the delay time if $d \geq 2^{32}$ (more than a few minutes);
 - at most 40% of the delay time if $d \geq 2^{28}$ (all other cases);
- **permissive licensing** – the prover software is open-sourced under the Apache 2.0 license.

3 VDF definitions

We start with recalling and adapting definitions of verifiable delay function protocol.

3.1 Generic VDF

Following the definition of VDF by [BBBF18] we consider only protocols of the following form:

- The protocol is executed by calling some function F , which takes a long time;
- After F finishes, one has to run another function Π to produce a delay proof π .

We say that the protocol (F, Π) is a *verifiable delay function* if it has the following features:

- **(Completeness)** Honest parties can execute it by calling F in t sequential steps and then by creating a proof π that will be accepted by a verifier. It is ok if there is a non-zero, but from all points of view negligible, chance to fail to produce a proof, even though many practical VDF proposals do provide *perfect completeness*, i.e. a proof can be always found.
- **(Fast verification)** The proof π can be verified in time polylogarithmic of t (i.e. exponentially faster than the time needed for computation). Note that this requirement rules out the designs based on computing an iterative function over a graph of a [special form](#), so that only a fraction of edges of the graph have to be proven, as these fractions turn to be bigger than polylogarithmic functions for practical parameters.
- **(Non-malleability)** For any starting point I and proof π it should be difficult to find ($O' \neq O = F(I)$) that can be verified by π . Here by difficulty we mean only the complexity of the online phase, and may ignore (or take with a smaller factor) the cost of precomputations.

- **(Sequentiality) or (Delay)** For any starting point I it should be difficult to compute (O, π) faster than in σt time for some σ close to 1. Note that we can allow asymptotically sublinear methods of computing as long as the constants are prohibitively large to be an actual time improvement.

It turns out that if we include quantum adversaries, then the only existing candidate is a combination of sequential algebraic transformation and a quantum-secure verifiable computation protocol. In order to boost the proof phase, the transformation F is chosen to be invertible with low algebraic degree of the inverse. The entire construction is then called *invertible VDF*.

3.2 VDF proofs based on Verifiable Computation

Verifiable Computation (VC) schemes are a family of proof of knowledge protocols, where Prover convinces Verifier that O is the result of program P applied to some input I , so that it is impossible to craft a proof that would pass the verification without actually computing O (= knowing all intermediate steps of computation). We see that this is exactly what is required for a delay proof, however, the existing VC schemes vary greatly in performance and requirements:

- **Scalability:** Prover time can be linear or superlinear of actual execution time. Note that our requirements potentially allow for superlinear proofs as long as they can be parallelized enough to make the total latency only moderately smaller than the execution time (Section 2.3);
- **Setup:** A proof system may utilize a so called common/structured reference string (CRS/SRS) as a source of entropy and hidden secrets. As the creator of SRS is able to forge proofs, one has to run complicated multi-party ceremonies in order to share the secret among participants and make no single party able to attack. Note that the size of CRS is usually proportional to the size of the computation (expressed as a circuit), so it can be very large for VDF. This makes us consider proof systems that do not require trusted setups [AHIV17, WTS⁺18, BBB⁺18, BBHR19]
- **Verifier scalability:** with a few exceptions [BBB⁺18], the VC proof systems allow verifier to be sublinear, if not constant-time, of the computational time. For the application of VDF constant-time and logarithmic-time verifiers are allowed, but a lot depends on the constant factor in the latter case.
- **Succinctness:** Proof size can be constant-size, logarithmic of program length, or even larger. The constant-size proofs are usually possible only in proof systems with trusted setup, which we have ruled out. As a result, logarithmic-size proofs seem to be the only option.

We expect that the delay function F will be tailored to the specific choice of the VC proof system.

3.3 Invertible VDF

In order to have delay granularity, it is convenient to have $F = f^D$ for some function f (called *base function*) and integer D being a multiple of the granularity parameter (see also). The reason is that the VC proof systems typically define computation as an arithmetic circuit with the statement of correctness being equivalent to a system of algebraic equations that relate intermediate steps of the computation. The cost of constructing a VC proof for VDF is determined by the degree of these equations, which is not always equal to the degree of f . Vice versa, we would be interested in function f such that the degree of its inverse would be much smaller as this would reduce proportionally the time needed to compute the proof π .

Definition 1. Let \hat{f} be the inverse of f . We say that f is γ -invertible if the cost of \hat{f} is γ times the cost of f . One-way functions over Z_N are at most N -invertible.

Remark 1. By using invertible functions for F we rule out the Hellman time-memory tradeoff [Hel80]. The latter is a generic method to find a preimage for function $F(x) \rightarrow y$ assuming that the attacker can precompute F on its entire domain (size of N) and store a certain part of this precomputation in the memory. The Hellman method tells us how to store these calls in the most efficient way so that a preimage to arbitrary y can be found in less than N calls. Clearly, this approach is not useful for invertible functions.

3.4 Invertible VDF flow

Section 3.1 provides a generic description of a VDF, but leaves out a couple of subtle details. Concretely, we want to address the typical scenario that prover is not allowed to start the protocol at arbitrary moment, but rather he is given a seed or a family of possible seeds. In fact, the latter is more probable and we want to emphasize the fact that prover has enough freedom/entropy in selecting the seed so that he can even guess a valid seed with non-negligible probability. We capture this scenario by using an input predicate ϕ , which is communicated to prover at the start of the protocol. However, the prover knows beforehand the fraction of possible inputs that ϕ will allow. For example, if a concrete application requires that invertible VDF is computed over $S_1 || S_2$ where S_1 is given to the prover and S_2 is arbitrary then $\phi(I) = \text{true}$ iff I starts with S_1 . Such inputs are easy to enumerate. A more elaborate scenario is that the VDF is computed over $H(S_1 || S_2)$ where H is a preimage-resistant hash function. In the latter case a Prover who does not know S_1 can't efficiently categorize his precomputations.

We thus suggest the following description:

1. **Setup** – Prover learns the input domain \mathcal{W} and input density, i.e. the (approximate) fraction of elements \mathcal{W} that can be selected as a seed.
2. **Start** – Prover is given *input predicate* ϕ over \mathcal{W} .
3. **Computation** – Prover selects I such that $\phi(I) = \text{true}$. Then he computes $O = F(I)$.
4. **Proof** – if $\chi(O) = \text{true}$, Prover creates an IVC proof π and submits (I, O, π) to Verifier.
5. **Verification** – Verifier checks π .
6. **Verification** – Verifier checks π .

Note that ϕ is crucial, otherwise it is possible to pre-select a valid O , invert F to some preimage I , and use it in further proofs.

Note that Prover can start computing the proof right after ϕ is known. However, he may spend some time selecting I if this reduces the computational time or cost.

Definition 2. We say that the protocol allows grinding if Prover is able to efficiently enumerate many I such that $\phi(I) = \text{true}$.

An example is when I must be a hash of some preimage I' , with the hash function being finally determined only in the setup phase.

3.5 Sequentiality

We would like to expand on the sequentiality property from Section 3.1 so that we capture precomputations. Indeed, if Prover is able to precompute F on the entire domain or a big fraction of it, and store the results in an efficiently accessible memory so that certain part of the **Computation** phase can be sped up, then the function is clearly unsuitable for a VDF.

Definition 3. The function F is (α, β, M) -sequential if M precomputation (calls to F) gives the chance α to reduce the running time of F to β .

In practice, we need that F is (α, β, M) -sequential for α and β close to 1 and M being the security level (in our case 128 bits).

As we consider security against quantum attacks, we should have an analog of sequentiality property for quantum machines.

Definition 4. *The function F is (α, β, M) -quantum-sequential if M precomputation (quantum calls to F) gives the chance α to reduce the running time of F to β .*

4 Attacks on invertible VDFs

In this section we explore generic attack strategies that are applicable to all instances of invertible VDFs that we consider.

4.1 Precomputation attacks

Here we consider attacks on sequentiality that arise from the ability of the attacker to precompute certain calls to F . Note that for the security level of 128 bits we allow an adversary to make 2^{128} precomputation calls to F . Recall that the VDF function F maps \mathcal{W} of size N to itself.

4.1.1 Attack 1

The simplest attack works as follows:

1. Apply F^{-1} (i.e. invert F) to the entire domain \mathcal{W} and create a proof for every inversion.
2. Store results as $\overline{\mathcal{W}} = \{(I, O, \pi)\}$.
3. When ϕ is known, find any $I \in \overline{\mathcal{W}}$ for which $\phi(I) = \text{true}$ and retrieve (O, π) from $\overline{\mathcal{W}}$.

This costs at most γN time and N space. Thus any function is $(1, 0, N)$ -sequential.

4.1.2 Attack 2

The same attack but available to the adversary running the equivalent of only M calls to F , works as follows:

1. Apply F^{-1} (i.e. invert F) to M/γ elements of \mathcal{W} and create a proof for every inversion.
2. Store results as $\overline{\mathcal{W}} = \{(I, O, \pi)\}$.
3. When ϕ is known, check if there is I such that $(I, *, *) \in \overline{\mathcal{W}}$ and $\phi(I) = \text{true}$. If found then retrieve (O, π) from $\overline{\mathcal{W}}$.

This costs at most M time and M/γ space, but the chance to find the right I is at most $\frac{M|\phi|}{\gamma N}$ with $|\phi|$ being the number of I for which $\phi(I)$ is true, so that the function is $(\frac{M|\phi|}{\gamma N}, 0, M)$ -sequential. If ϕ is chosen so that it is not trivial to find the intersection of $\overline{\mathcal{W}}$ and $|\phi|$, and one has to check each entry of $\overline{\mathcal{W}}$ separately, then the chance to find the right I is only $\frac{GM}{\gamma N}$ where G is the number of calls to ϕ that are made by the adversary (possibly in parallel) during the time required by application to compute VDF. In this scenario the VDF is at best is

$$(\frac{GM}{\gamma N}, 0, M)\text{-sequential.}$$

Example 1: suppose $M = 2^{64}$, $G = 2^{58}$, $N = 2^{128}$, $\gamma = 1/64$. Then an adversary who spends time equivalent to 2^{64} VDF forward precomputations, can compute the F output after trying 2^{58} values for ϕ .

4.1.3 Attack 3

Another modification of the same attack¹ works as follows:

1. Compute not the full $F^{-1} = f^{-D}$ but the reduced version $f^{\delta D}$, i.e. compute chains of length δD for some $\delta > 0$ (and parts of future proofs if possible). In total it is possible to compute $\frac{M}{\gamma\delta}$ chains.
2. Store results as $\overline{W} = \{(I', O')\}$.
3. When ϕ is known, select up to G different values I such that $\phi(I) = \text{true}$ and start computing. For each intermediate call to f check if the state matches any of \overline{W} . If found then use the entry to reduce the computing time by the factor δ .

The chance to meet a precomputed chain in the online phase is $\frac{(1-\delta)DGM}{\delta\gamma N}$, so F is at best

$$\left(\frac{(1-\delta)DGM}{\delta\gamma N}, 1-\delta, M\right)\text{-sequential.}$$

Note that this attack equally works for F where each f call is different. In this case each precomputation chain simply starts from a valid output state.

Example 2: suppose $D = 2^{40}$, $M = 2^{64}$, $N = 2^{128}$, $\gamma = 1/64$, $\delta = 1/2$. Then an adversary who spends time equivalent to 2^{64} VDF forward precomputations, can compute VDF in $1/2$ of required time if he can compute 2^{18} VDF chains in parallel.

4.1.4 Mitigation and parameter selection

Looking at the sequentiality parameters, it is easy to see that the attacks get non-negligible probability to succeed whenever GM approaches N/D where G is the amount of computation spent in the online phase, M is the amount of precomputation, and D is the VDF difficulty parameter. Therefore it is natural to increase N to mitigate it, for example to make $N = M^{1.5}$ (a bit risky) or $N = M^2$ or even bigger. Assuming $G < M$ we can state that the security level of invertible VDF is roughly given by the value $\sqrt{N/D}$:

$$\text{Security level (VDF)} \approx \sqrt{\frac{N}{D}} \quad (1)$$

Thus, in order to have 100 bits of security we need a 256-bit long N .

It is also clear that a simple ϕ also makes the attack easier. Thus we suggest using a preimage-resistant hash function as part of ϕ , but not part of the proof. There can be many possible such ϕ , and we suggest the already mentioned

$$\phi_{S_1}(I) = \text{true} \Leftrightarrow \exists S_2 : H(S_1 || S_2) = I \quad (2)$$

where H is a regular cryptographic hash function (such as SHA-3) and S_1 is the parameter selected at the **Start** phase.

The ability to check sufficiently many different S_2 is also called *grinding*. In the further analysis we always assume that it is possible to grind up to 2^{64} different inputs; if this is not the case in some application, the protocol is even more secure.

4.2 Grobner Basis and Algebraic

This class of attacks exploits the algebraic nature of an invertible VDF. Consider the reduced F as $F' = f^d$ for some $d < D$. If it is possible to solve the equation

$$\hat{f}^d(I) = O \quad (3)$$

¹Suggested to us by Avihu Levi

where O is known and X is unknown, then one can run the equation solver instead of applying f^d . This attack only succeeds if (3) can be solved fast then running f^d .

If (3) is multivariate, then the best generic technique is applying Groebner basis algorithms [CLO97, Fau02]. These algorithms find a generating set for the ideal generated by (3) so that they can be solved as univariate equations. Little is known about the actual complexity of Groebner basis algorithms, but the available heuristic analysis suggests that it is exponential in d , thus quickly become impractical. Even in the unlucky case when a Groebner basis can be found for the system of equations that describe a VDF, we still face the problem of solving univariate equations for it.

Univariate equations over finite fields can be solved using Berlekamp or Cantor-Zassenhaus algorithms. It is known [CZ81] that an equation of degree r can be solved using $O(r^3 \log p)$ operations over elements of \mathbb{F}_p . This should be compared with $O(d \log p)$ operations needed to compute f^d . We conclude that, unless f^d has some structure that makes the univariate equations of its Groebner basis of very low degree (at most $\sqrt[3]{d}$), generic algebraic methods are not applicable even if f or \hat{f} has univariate representation.

5 Quantum algorithms

Quantum algorithms are the algorithms specially designed for quantum computers. A unique property of a quantum computer is able to perform operations not on a single state but on a superposition of states, so that certain classical algorithms can be amortized. The most famous examples are Shor's algorithm for factorization [Sho94] and its derivative for discrete logarithm search, as well as the Grover search algorithm [Gro96]. The latter is our main interest as we have excluded factoring- and DLP-based schemes from consideration.

5.1 Grover's algorithm

The Grover algorithm was originally formulated as a database search, however, it can be equivalently written as a preimage search algorithm for some hard-to-invert function. Concretely, given $O \in Z_N$ (or any other domain efficiently mappable to Z_N) and function H , Grover algorithm finds I such that $H(I) = O$ in time equivalent to \sqrt{N} calls to H on a quantum computer. Note that the algorithm is inherently sequential and so far no one knows how to parallelize it so that the computing time can be reduced. The probability to find the preimage in t time grows as t^2/N [BHT97].

5.1.1 Single-target Grover's

Grover's algorithm makes a particular sequence of operations multiple times. At a high-level, after each iteration the probability amplitude of the marked index state is increased. Once enough iterations are performed, this amplitude is high enough that when a measurement is performed, with high probability the resulting index corresponds to the special marked state. Each iteration is made of 2 simple operations. The following description assumes that we are using Grover's algorithm to perform unstructured search on a boolean function $f : \{0, 1\}^n \rightarrow \{0, 1\}$ where there is only one unique element x' such that $f(x') = 1$ and $N = 2^n$. The input to Grover's algorithm is an equal superposition of all inputs:

$$\phi_0 = \frac{1}{\sqrt{N}} \sum_{x=0}^{N-1} |x\rangle$$

The 2 operations performed at each iteration of Grover's algorithm are as follows:

- **Amplitude flip** $|x'\rangle \rightarrow -|x'\rangle$: The amplitude of the unique element x' is negated. This operation is implemented using the following gate logic:

$$(-1)^{f(x)} |x\rangle = \begin{cases} |x\rangle & \text{if } f(x) = 0 \\ -|x\rangle & \text{if } f(x) = 1 \end{cases}$$

Post this operation, the state of the input qubits is:

$$-\frac{1}{\sqrt{N}} |x'\rangle + \sum_{x \neq x'} \frac{1}{\sqrt{N}} |x\rangle$$

- **Diffusion operator:** The goal of the diffusion operator is to increase the amplitude of x' . The diffusion operator flips the amplitude of each $x \in \{0,1\}^n$ around the average amplitude $\mu = \frac{1}{N} \sum_x A_x$ where A_x represents the amplitude of x . This operation is implemented using the following gate logic:

$$\sum_{x \in \{0,1\}} \alpha_x |x\rangle \mapsto \sum_{x \in \{0,1\}^n} (2\mu - A_x) |x\rangle$$

The mean amplitude before the diffusion operator is roughly $\frac{1}{\sqrt{N}}$. Therefore, after the diffusion operator all entries remain roughly the same as they go from, $\frac{1}{\sqrt{N}}$ to $\frac{2}{\sqrt{N}} - \frac{1}{\sqrt{N}}$. However, the amplitude of x' , $A_{x'}$, goes from $-\frac{1}{\sqrt{N}}$ to $\frac{3}{\sqrt{N}}$.

The above two operations are then repeatedly applied, increasing the amplitude of x' after each round by $\frac{2}{\sqrt{N}}$. After $O(\sqrt{N})$ rounds, the amplitude of x' is up to a desirable constant and at this point a measurement is performed which collapses the state to x' with constant probability.

5.1.2 Multi-target Grover's

For us the main interest is the multi-target version of the Grover algorithm, which we hope to use in the pre-computation attacks. The same algorithm works in the multi-target case but now instead of requiring $O(\sqrt{N})$ iterations, the number of iterations required is $O(\sqrt{\frac{N}{k}})$ as there are k targets whose amplitude gets amplified after each iteration. Hence, the multi-target Grover's finds a pre-image to one of k targets in time $\sqrt{N/k}$. The probability to find the pre-image in t time grows as

$$f_{MG}(k, t, N) = kt^2/N \tag{4}$$

according to [BHMT02].

5.1.3 Quantum Parallelism

Achieving speed-ups via parallelism is an important line of work for most search algorithms. However, in the case of Grover's algorithm, the work of Zalka [Zal99] proves that quantum searching can not be parallelized better than the naive case. The naive case assigns different parts of the search space to separate quantum computers. And hence, you get a linear improvement in search time with multiple quantum computers corresponding to increased computational cost.

5.2 Attacks based on Grover's

5.2.1 Grover Attacks 1 and 2

We can apply Grover algorithm to our precomputation attacks from Section 4.1. The first two attacks work as follows:

1. Apply F^{-1} (i.e. invert F) to M/γ elements of \mathcal{W} and create a proof for every inversion.

2. Store results as $\overline{W} = \{(I, O, \pi)\}$.
3. When ϕ is known, check (using Grover) if there is I such that $(I, *, *) \in \overline{W}$ and $\phi(I) = \text{true}$. If found then retrieve (O, π) from \overline{W} .

The improvement comes in the setting when finding the right I is not easy i.e. when we have to check each $\phi(I) = \text{true}$ individually. We run Grover algorithm on the set \overline{W} of targets in the Computation phase, and in order to increase the attack probability we run several (denoted by G) copies of Grover in parallel. Let us assume that a call to ϕ costs as much as a call to f . Then if each Grover is run for βD , then the chance to match any element of \overline{W} can be derived from (4) as

$$P_{MG}(\beta, D, N) = \frac{\beta^2 G D^2 M}{\gamma N} \quad (5)$$

so that the VDF protocol is

$$\left(\frac{\beta^2 G D^2 M}{\gamma N}, \beta, M \right)\text{-quantum-sequential}$$

Example 3: suppose $M = 2^{64}$, $D = 2^{40}$, $N = 2^{192}$, $G = 2^{40}$, $\gamma = 1/64$. Then an adversary who precomputes an equivalent of 2^{64} VDF forward precomputations, can compute the F output in $1/2$ time using 2^{40} instances of Grover algorithm with probability $1/4$.

5.2.2 Attacks that do not work

Interestingly, Grover algorithm does not seem to improve if we switch from attacks 1-2 to the chain-based approach (Section 4.1.3). The reason is that the latter requires to compute several iterations of f which can not be iterated in Grover. Thus we conclude that (5) is the best we can hope with Grover.

For the same reason we can not effectively run a quantum meet-in-the-middle attack [BHT97] on F by splitting in two parts, as Grover would need to compute many iterations of F sequentially.

We conclude that the security level of invertible VDFs against quantum adversaries is given by the formula

$$\text{Security level (VDF)} \approx \frac{\sqrt{N}}{D} \quad (6)$$

which is only by the factor of \sqrt{D} lower than the classical one (1).

5.3 Practicality of Quantum Search

Grover's search is usually discussed for some oracle function $f(\cdot)$. However, several practical issues need to be considered when a specific construction of $f(\cdot)$ is being considered. The work of Viamontes, Markov and Hayes [VMH05] focuses on these issues. The first issue is that in order to query $f(\cdot)$ in superposition, we need a quantum hardware implementation of $f(\cdot)$. While a quantum circuit is usually similar in size to its classical counterpart, if the circuit's maximum depth exceeds \sqrt{N} then the evaluation of $f(\cdot)$ is the more cost consuming part of searching a size N database. This would diminish the quantum speedup by a considerable amount.

5.4 Quantum Grobner

We note that quite little is known about the quantum complexity of Groebner basis algorithms nor of the equation solving. There exist several results showing that Gaussian elimination of sparse matrices can be sped up greatly on a quantum computer <https://arxiv.org/abs/0811.3171>, which in turns would shorten the elimination step in the Groebner basis algorithms. However, it is the subject of the future research if an algorithm like F4 or F5 [Fau02] can be improved as a whole. We expect that if it can be done, running a quantum algorithm every few steps would induce high latency, and the high value of D would rule out such strategies. To the best of our knowledge there is no quantum improvement to solving univariate equations either.

6 Other attacks

6.1 MiMC reduction

Here we provide an observation that relates the sequentiality of invertible VDFs to the security of some well-known symmetric key primitives. For this we recall the MiMC bijective transformation over \mathbb{F}_p^2 [AGR⁺16]:

- Input $(x, y) \in \mathbb{F}_p^2$
- For $1 \leq i \leq R$ repeat:
 1. $(x, y) \leftarrow (y + x^3 + c_i, x)$ where c_i is a public constant that depends on i . In a keyed version, a key is also added mod p with c_i .

The security of MiMC is based on the assumption that the resulting polynomial has large algebraic degree and it has no compact representation. Despite the years of public scrutiny [ACG⁺19, EGL⁺20], no counterargument to this assumption was found². Computing Groebner basis has proven useless either. This hints that as long as the function f in the VDF design resembles that of MiMC, we should expect little to no algebraic structure in the polynomial describing the iteration of multiple rounds, which in turn implies that these polynomials are kind of “generic” and “random-looking”. Thus intuitively there should be no sequentiality shortcut as it should not be for random functions.

We note that this argument holds as long as the polynomial describing the iteration has degree 3. It is known that certain quadratic polynomials yield short cycles and thus should not be used [VS04].

6.2 Attacks on Verifiable Computation

In this report we mostly ignore potential attacks on the proof system used in verifiable computation systems. We only note that they largely split into two groups: ones whose security is based on the hardness of discrete logarithm in elliptic curve or integer groups (and thus almost universally quantum insecure), and the ones that require little to no such assumption. However, not all protocols from the latter group have claimed post-quantum security as some of their components such as Fiat-Shamir challenge derivation are not proven post-quantum secure. We thus urge that in order to achieve post-quantum security it is desirable to use only those VC systems that do claim it (even though no formal proof of security may be given).

7 Concrete Proposals

All proposals are considered in two settings: 128-bit field and 256-bit field, commonly denoted \mathbb{F}_p .

7.1 Recommendations for ϕ

One can note that some VDF proposals employ a hash function at the first step of the computation, i.e. as part of ϕ in our terminology, whereas others do not. We have explained in Section 4.1.4 that the use of hash function makes all precomputation attacks universally harder, whereas it costs little for prover and verifier. We thus include it into our proposal (MinRoot) and advise all others do the same.

²There are certain cryptanalytic results for the binary field case and in the related key scenario, but they are irrelevant for our purpose.

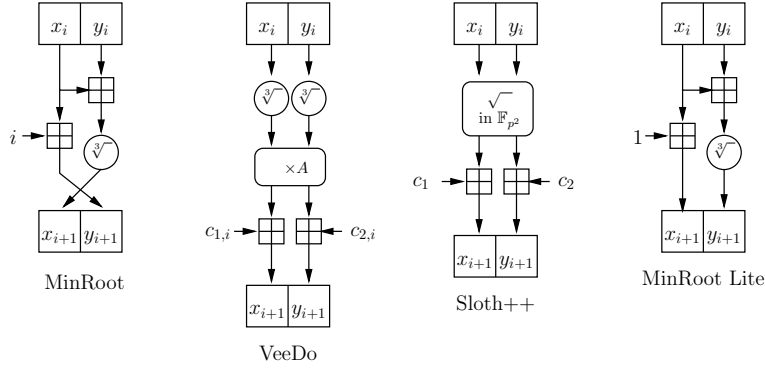


Figure 1: One round of MinRoot and MinRoot Lite

7.2 MinRoot (Khovratovich-Feist)

For our proposal MinRoot we select function f as part of F so that the inverse \hat{f} is fastest possible to compute in the proof creation part. Concretely, we consider function $h(x) = x^{(2p-1)/3}$, which is the inverse of a cubic (i.e. low degree) function: $h^3(x) = x^{2p-1} = x \bmod p$. Then we define

1. Input predicate ϕ : x_1 is the seed, input is $I = (x_1, y_1)$, and the relation is $\phi(x_1, y_1) = \text{true} \leftrightarrow y_1 = H(x_1)$ where H is SHA-3.
2. Computation: repeat for $1 \leq i \leq D$

$$(x_{i+1}, y_{i+1}) \leftarrow ((x_i + y_i)^{(2p-1)/3}, x_i) + (0, i)$$

3. Output $O = (x_{D+1}, y_{D+1})$.
4. Proof: create a VC proof (presumably STARK-based) that $F(I) = O$.

We aim for the performance here: at each step we make one exponentiation, and those are sequential.

MinRoot claims 128-bit security for 256-bit primes.

7.3 VeeDo (StarkWare)

The VeeDo function is also based on the cubic roots in \mathbb{F}_p :

1. Input predicate ϕ : x_1 is the seed, input is $I = (x_1, y_1)$, and the relation is $\phi(x_1, y_1) = \text{true} \leftrightarrow y_1 = x_1$.
2. Computation: repeat for $1 \leq i \leq D$

$$(x_{i+1}, y_{i+1}) \leftarrow A \cdot (x_i^{(2p-1)/3}, y_i^{(2p-1)/3}) + (c_{1,i}, c_{2,i})$$

where A is an invertible matrix and $\{c_{1,*}, c_{2,*}\}$ are some public constants.

3. Output $O = (x_{D+1}, y_{D+1})$.
4. Proof: create a VC proof (presumably STARK-based [BBHR19]) that $F(I) = O$.

The difference to MinRoot is that a combination of a power function and the state swap is replaced with a full-blown affine transformation. This approach, well known in the theory of block cipher design, is supposed to provide faster diffusion and to ease analysis. However, we note that the diffusion speed is not as important for our setting as sequentiality is, since after 2^{20} or more iterations the diffusion becomes almost perfect anyway. Regarding sequentiality, we remark that VeeDo spends two parallelizable exponentiations per round, whereas MinRoot does only one. Thus for the same delay time and same state size VeeDo requires twice bigger hardware, whereas the smaller state would reduce the security. The same ratio holds for the proof size and computation time: VeeDo requires twice as much for the a We thus conclude that VeeDo is slightly inferior to MinRoot

7.4 Sloth and Sloth++ (Wesolowski-Boneh et al.)

Sloth was suggested in [LW17]. For the k -bit a security level we select a prime p so that $p = 4q + 3$ for some q and $p > 2^{2k}$. Let H be a regular cryptographic hash function such as SHA-2 with $2k$ -bit output. The algorithm is as follows:

1. Input predicate ϕ : x_1 is the seed, input is $I = (x_1, y_1)$, and the relation is $\phi(x_1, y_1) = \text{true} \Leftrightarrow y_1 = H(x_1)$.
2. Computation: repeat for $1 \leq i \leq D$

$$y_{i+1} \leftarrow (\delta \sigma(y_i))^{(p+1)/4} \in \mathbb{F}_p$$

where σ is a not fully specified “simple” invertible function such as bit permutation, and $\delta \in \{1, -1\}$ is chosen so that the result is quadratic residue.

3. Output $O = y_{D+1}$.
4. Proof: direct evaluation of the inverse function.

Boneh et al suggested in [BBBF18] a version more suitable for verifiable computation proofs. It is called Sloth++ and it operates similarly over \mathbb{F}_{p^2} :

1. Input predicate ϕ : s is the seed, input is $I = (s, x_1, y_1)$, and the relation is $\phi(s, x_1, y_1) = \text{true} \Leftrightarrow (x_1, y_1) = H(s)$.
2. Computation: repeat for $1 \leq i \leq D$

$$(x_{i+1}, y_{i+1}) \leftarrow \underbrace{(x_i, y_i)^{(p^2+1)/4}}_{\text{as element of } \mathbb{F}_{p^2}} + (c_1, c_2)$$

3. Output $O = (x_{D+1}, y_{D+1})$.
4. Proof: some VC proof.

In order to compare apples to apples, we consider both 128- and 256-bit primes for Sloth++, i.e. the same as for VeeDo and MinRoot. For l -bit prime we get that one exponentiation to an l -bit exponent (MinRoot) or two (parallelizable) exponentiations (VeeDo) are replaced with a single $2l$ -bit exponentiation of a $2l$ -bit field element. This results in approximately twice longer rounds. Thus from the performance perspective the Sloth++ based on 128-bit prime is similar to MinRoot with 256-bit prime, and Sloth++ based on a 256-bit prime is similar to MinRoot with a 512-bit prime. Securitywise, it is unclear whether the algebraic structure of \mathbb{F}_{p^2} can be exploited in precomputation or quantum attacks, but we should not rule this option out. Note that the authors of Sloth++ claim only 64 bits of security for Sloth++ based on 128-bit primes, which is less than we (maybe optimistically) claim for MinRoot with 256-bit primes (see below).

7.5 MinRoot Lite

We also bring into a consideration a memory-optimized version of MinRoot, where we replace one variable with just a counter:

1. Input predicate ϕ : x_1 is the seed, input is $I = (x_1, y_1)$, and the relation is $\phi(x_1, y_1) = \text{true} \leftrightarrow y_1 = H(x_1)$ where H is SHA-3.
2. Computation: repeat for $1 \leq i \leq D$

$$(x_{i+1}, y_{i+1}) \leftarrow ((x_i + y_i)^{(2p-1)/3}, y_i) + (0, i)$$
3. Output $O = (x_{D+1})$.
4. Proof: create a VC proof (presumably STARK-based [BBHR19]) that $F(I) = O$.

The resulting state is twice smaller thus precomputation attacks are more efficient. However, the memory footprint is minimal as we do not have to store almost anything except for the number we exponentiate. This makes it similar to Sloth++.

8 Concrete security of MinRoot(s)

Recall that the function f in MinRoot is defined as

$$f(x, y) = ((x + y)^{(2p-1)/3}, x + i), \quad x, y \in \mathbb{F}_p$$

where p is 128- or 256-bit and i is the round number. The number of multiplications needed to compute f is the length of *optimal addition chain* of $(2p-1)/3$, denoted by $l((2p-1)/3)$. It is [known](#) that

$$l(n) > \log_2 n + \text{wt}(n) - 3$$

which is about $1.5 \log_2 p$ in our case. Since the inverse function \hat{f} needs 2 multiplications, we obtain

$$\gamma = \frac{\text{Multiplications for } \hat{f}}{\text{Multiplications for } f} \approx \begin{cases} 2^{-6.5} & \text{for 128-bit } p; \\ 2^{-7.5} & \text{for 256-bit } p \end{cases}$$

Note that in the context of precomputation and other attacks the state has effective size of $2 \log p$ bits.

8.1 Classical attacks

First we consider attacks from Section 4.1. In attack 2 the function is $(\frac{GM}{\gamma N}, 0, M)$ -sequential. Assuming the adversary spends as much as 2^{128} precomputations, we get that the VDF is

$$(\frac{G}{2^{121.5}}, 0, 2^{128})\text{-sequential for 128-bit } p; \tag{7}$$

$$(\frac{G}{2^{248.5}}, 0, 2^{256})\text{-sequential for 256-bit } p. \tag{8}$$

In attack 3 the function is $(\frac{(1-\delta)DGM}{\delta \gamma N}, 1 - \delta, M)$ -sequential. Again, assuming the adversary spends as much as 2^{128} precomputations and the difficulty parameter is the most favourite $D = 2^{44}$, we get for $\delta = 0.5$ that the VDF is

$$(\frac{G}{2^{78}}, 0.5, 2^{128})\text{-sequential for 128-bit } p; \tag{9}$$

$$(\frac{G}{2^{269}}, 0.5, 2^{192})\text{-sequential for 256-bit } p. \tag{10}$$

Thus we claim 80 bits of security for 128-bit primes and 192 bits of security for 256-bit primes. If the precomputation capacity of the adversary does not exceed his computation phase costs, then we can claim 100 bits of security for 128-bit primes.

8.2 Quantum attacks

Recall that the best quantum attack we have found combines the multi-target Grover with precomputations and yields the VDF function being

$$\left(\frac{\beta^2 G D^2 M}{\gamma N}, \beta, M\right)\text{-quantum-sequential}$$

Again, assuming the adversary spends as much as 2^{128} (classic) precomputations and the difficulty parameter is the most favourite $D = 2^{44}$, we get for $\beta = 0.5$ that the VDF is

$$\left(\frac{G}{2^{36}}, 0.5, 2^{128}\right)\text{-quantum-sequential for 128-bit } p; \quad (11)$$

$$\left(\frac{G}{2^{84}}, 0.5, 2^{80}\right)\text{-quantum-sequential for 128-bit } p; \quad (12)$$

$$\left(\frac{G}{2^{228}}, 0, 2^{192}\right)\text{-sequential for 256-bit } p. \quad (13)$$

Thus it has at most 36 bits of quantum security for 128-bit p with 2^{128} precomputations of F , but 80 bits of quantum security for 128-bit p if only 2^{80} precomputations of inverted VDF are allowed. Note that the latter setting is actually quite demanding for the adversary, as each inverted VDF is 2^{46} exponentiations, so that the 2^{80} inverted VDFs correspond to about 2^{120} calls to AES or SHA-3.

9 Concrete security of other proposals

Our analysis for MinRoot applies equally to VeeDo and Sloth++ with the same prime size. However, the γ value differs slightly: we spend $3 \log p$ operations in the forward direction and at least 8 multiplications in the backward direction as we apply the inverse of A . Therefore, we have $\gamma = 2^{-6}$ for 128-bit p and 2^{-7} for 256-bit p .

The security of MinRoot Lite is lower: due to smaller state the security of MinRoot Lite with 256-bit primes corresponds to the security of MinRoot with 128-bit primes: 80-bit classic and 40-bit quantum security with 2^{128} VDF precomputations, 100-bit classic and 64-bit quantum with 2^{100} VDF precomputations.

10 Performance

We have not fully benchmarked the proposals but since all of them use pretty much the same field operations, we can still reliably compare them by counting the number of field operations. A multiplication in \mathbb{F}_{p^2} is assumed equivalent to 4 multiplications in \mathbb{F}_p . Thus a single call to f , w.l.o.g., a 512-bit state (2 256-bit prime field variables) requires

- 300 multiplications in MinRoot and MinRoot Lite;
- 600 multiplications in VeeDo;
- 600 multiplications in Sloth++.

We estimate a single multiplication in a 128-bit field to take 0.5 ns on dedicated hardware, so 1 round of MinRoot to take about 75 ns, and 300ns for 256-bit modulus. Regular CPUs seem to bring significant overhead, as [reports](#) 10K CPU cycles per square root computation over 256-bit modulus

11 Conclusion

We conclude that MinRoot has security level equal to that of VeeDo and Sloth++, but slightly exceeds these designs in performance (spending less computations for the same sequentiality settings) and in the design simplicity. We thus recommend MinRoot for VC-based VDF.

References

- [ACG⁺19] Martin R. Albrecht, Carlos Cid, Lorenzo Grassi, Dmitry Khovratovich, Reinhard Lüftenegger, Christian Rechberger, and Markus Schofnegger. Algebraic cryptanalysis of stark-friendly designs: Application to marvellous and mimc. In *ASIACRYPT (3)*, volume 11923 of *Lecture Notes in Computer Science*, pages 371–397. Springer, 2019.
- [AGR⁺16] Martin R. Albrecht, Lorenzo Grassi, Christian Rechberger, Arnab Roy, and Tyge Tiessen. Mimc: Efficient encryption and cryptographic hashing with minimal multiplicative complexity. In *ASIACRYPT (1)*, volume 10031 of *Lecture Notes in Computer Science*, pages 191–219, 2016.
- [AHIV17] Scott Ames, Carmit Hazay, Yuval Ishai, and Muthuramakrishnan Venkitasubramaniam. Liger: Lightweight sublinear arguments without a trusted setup. In *ACM Conference on Computer and Communications Security*, pages 2087–2104. ACM, 2017.
- [BBB⁺18] Benedikt Bünz, Jonathan Bootle, Dan Boneh, Andrew Poelstra, Pieter Wuille, and Greg Maxwell. Bulletproofs: Short proofs for confidential transactions and more. In *IEEE Symposium on Security and Privacy*, pages 315–334. IEEE, 2018.
- [BBBF18] Dan Boneh, Joseph Bonneau, Benedikt Bünz, and Ben Fisch. Verifiable delay functions. In *CRYPTO (1)*, volume 10991 of *Lecture Notes in Computer Science*, pages 757–788. Springer, 2018.
- [BBHR19] Eli Ben-Sasson, Iddo Bentov, Yinon Horesh, and Michael Riabzev. Scalable zero knowledge with no trusted setup. In *CRYPTO (3)*, volume 11694 of *Lecture Notes in Computer Science*, pages 701–732. Springer, 2019.
- [BHMT02] Gilles Brassard, Peter Høyer, Michele Mosca, and Alain Tapp. Quantum amplitude amplification and estimation. *Quantum Computation and Information*, page 53–74, 2002.
- [BHT97] Gilles Brassard, Peter Høyer, and Alain Tapp. Quantum cryptanalysis of hash and claw-free functions. *SIGACT News*, 28(2):14–19, 1997.
- [CLO97] David A. Cox, John Little, and Donal O’Shea. *Ideals, varieties, and algorithms - an introduction to computational algebraic geometry and commutative algebra (2. ed.)*. Undergraduate texts in mathematics. Springer, 1997.
- [CZ81] David G Cantor and Hans Zassenhaus. A new algorithm for factoring polynomials over finite fields. *Mathematics of Computation*, pages 587–592, 1981.
- [EGL⁺20] Maria Eichlseder, Lorenzo Grassi, Reinhard Lüftenegger, Morten Øygarden, Christian Rechberger, Markus Schofnegger, and Qingju Wang. An algebraic attack on ciphers with low-degree round functions: Application to full mimc. In *ASIACRYPT (1)*, volume 12491 of *Lecture Notes in Computer Science*, pages 477–506. Springer, 2020.
- [Fau02] Jean Charles Faugere. A new efficient algorithm for computing gröbner bases without reduction to zero (f 5). In *Proceedings of the 2002 international symposium on Symbolic and algebraic computation*, pages 75–83, 2002.

- [Gro96] Lov K. Grover. A fast quantum mechanical algorithm for database search. In *STOC*, pages 212–219. ACM, 1996.
- [Hel80] Martin E. Hellman. A cryptanalytic time-memory trade-off. *IEEE Trans. Inf. Theory*, 26(4):401–406, 1980.
- [LW17] Arjen K. Lenstra and Benjamin Wesolowski. Trustworthy public randomness with sloth, unicorn, and trx. *Int. J. Appl. Cryptogr.*, 3(4):330–343, 2017.
- [Pie19] Krzysztof Pietrzak. Simple verifiable delay functions. In *ITCS*, volume 124 of *LIPICs*, pages 60:1–60:15. Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 2019.
- [Sho94] Peter W. Shor. Algorithms for quantum computation: Discrete logarithms and factoring. In *FOCS*, pages 124–134. IEEE Computer Society, 1994.
- [VMH05] George F. Viamontes, Igor L. Markov, and John P. Hayes. Is quantum search practical? *Comput. Sci. Eng.*, 7(3):62–70, 2005.
- [VS04] Troy Vasiga and Jeffrey O. Shallit. On the iteration of certain quadratic maps over $\text{gf}(p)$. *Discret. Math.*, 277(1-3):219–240, 2004.
- [Wes19] Benjamin Wesolowski. Efficient verifiable delay functions. In *EUROCRYPT (3)*, volume 11478 of *Lecture Notes in Computer Science*, pages 379–407. Springer, 2019.
- [WTS⁺18] Riad S. Wahby, Ioanna Tzialla, Abhi Shelat, Justin Thaler, and Michael Walfish. Doubly-efficient zkSNARKs without trusted setup. In *IEEE Symposium on Security and Privacy*, pages 926–943. IEEE, 2018.
- [Zal99] Christof Zalka. Grover’s quantum searching algorithm is optimal. *Physical Review A*, 60(4), 1999.